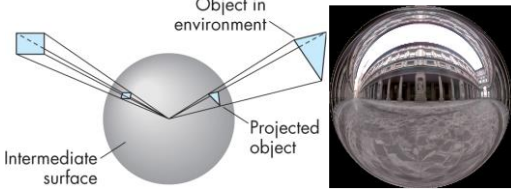


## Final Study Guide – Nathan Tung

### Lecture 7:

- Review: transparency order matters: draw over (render after) to show up on top (or flip for different blending)
- Environment Mapping (same as reflection mapping)
  - Reflection of a **fixed** environment; reflections computed via ray tracing (too slow in real-time environments)
  - Approximate the effect using texture map instead
    - Compute reflection vector (i.e. eye  $\rightarrow$  surface  $\rightarrow$  reflect  $\rightarrow$  world); intersect that ray with scene and compute color (shading, actually) value; unfortunately, this is pretty much ray tracing
  - Approximate with two-pass method:
    - Place camera at mirror object location, facing in direction of mirror surface normal
    - Render scene without mirror into texture map
    - Render scene normally with texture map applied to mirror
    - This time with original camera position, mirror back in scene
    - Difficulties: where to put the camera? mirror is missing in first render (messes up light/mirror occludes something), projection from mirror can be tricky/off-axis, and we have to re-render texture every time camera OR mirror moves
  - How can we do environment mapping? Project scene onto sphere at COP
    - Viewer can't tell difference between object vs. projection (ex. Planetarium)
    - Given reflection vector, we find shaded color in scene and determine texture coordinates (s,t)
    - Not correct, but good enough (created at specific location like origin, not location of mirrored surface)
      - Need to re-compute for scene changes
      - Difficult but doable to create a map (real cameras like Google Street View use spherical lenses)
    - OpenGL is simpler: use projections, or even simpler, cube mapping!
      - For cube, render 6 images, each centered on an axis
      - Results are inside of a cube, unfolded; FOV must be 90 deg with matching edges though
      - Obviously, need to pass correctly transformed surface normal into vertex shader
      - Can compute reflection vector, pass into fragment shader (gets interpolated in fragment shader)
      - Notice: samplerCube type (similar to sampler2D) and textureCube function
- Bump Mapping (really displacement mapping)
  - Texture map stores displacements to the surface; requires computing normal at the new, displaced surface point
  - This needs partial derivatives to be solved; finite differences can be solved; slow when solving every fragment
  - Advantage is to use the result to perturb the normal in object space
  - Normal mapping: what we typically think of as bump mapping, but we're storing normal in texture map
    - X, y, z components of a normal vector are stored in the RGB channels of texture map
    - Need to map [-1,1] range of normals to [0,1] range of colors; use  $[R,G,B] = ([x,y,z]+1)/2$
    - Not enough: still need a new "coordinate frame" and "Tangent Space" (or TBN space)
      - Lighting for surfaces with normal maps applied are computed in TBN space
      - Tangent: tangent to surface at point p
      - Bi-tangent: cross product of the normal and tangent
      - Normal: surface normal at point p
    - TBN vectors define new frame at point we are lighting on surface; align T and B vectors with s and t dimensions of normal map; get object space to tangent space by a vector transformation
    - $[T_x, T_y, T_z; B_x, B_y, B_z; N_x, N_y, N_z]$ ; transform light vector by this matrix
    - Light and normal map vectors are in same coordinate system, so lighting equations are correct

### Lecture 8:


- Picking and Selection
  - Use mouse to select item; several methods available
    - Basic ones: old glSelect() mechanism; color buffer picking
    - Advanced ones: occlusion query extension; ray casting
  - glSelect() in OpenGL since v1.0, deprecated after v3.0, basically gone now
    - Very slow (implemented in software, stalls hardware rendering pipeline)
    - Idea is to use integers to "name" objects in the scene and mark stream of geometry with them

- Works by rendering scene in special mode; buffer allocated to collect “hits”
      - A “hit” being any “name” intersecting view volume
      - gluPickMatrix() then helps restrict projection to area around mouse click (ex. 4x4 area of pixels)
      - Any primitives rendered after a “name” object defined is placed in predefined hit buffer
      - “Name” objects in hit buffer tell us what was under the mouse
  - Color Buffer Picking
    - Much simpler/faster; uses additional rendering pass
    - Every object we want to identify is given unique color (simple fragment shader, so directly assign color)
    - No lighting or texture mapping done; otherwise, scene is rendered normally
      - Z-buffer on, double buffering required (or rendered into dedicated color buffer for picking), only draw objects we wish to consider for picking
    - Once render pass is complete, DON’T swap buffers;
      - Call gl.readPixels() to recover result using mouse/window (x,y)
    - Pixel color value retrieved corresponds to object underneath the mouse
    - Variation of this technique uses separate color buffer rendered in parallel to normal buffer
    - Further enhancements for speeding it up:
      - Render only bounding volumes for scene objects
      - Be sure not to try/identify more objects than there are bits of precision in buffer
      - Only render objects that need to be selected
    - Problematic for objects with transparent textures (alpha values!), but manageable with some effort
  - Other techniques
    - Occlusion query – a special case
      - Want to know if something will be visible before we render it
      - Reading z-values instead of color
      - Allows us to determine order of objects under mouse
      - Harder to do region queries (e.g. rubber band box)
    - Ray Casting
      - Requires objects to be in memory (in some form)
      - Usually bounding solids used
      - Does not work at pixel level
      - Can have trouble with transparency (e.g. picking a hole)
- Collision Detection
  - Object Representation and Bounding Volumes, Simple Intersection Tests, Bounding Volume Hierarchies
    - Has object A collided with object B?
  - Object Representation
    - Compute whether one primitive intersects another
    - Could be done by comparing all triangles in a scene with one another, but slow/produces useless results
    - Stick only to objects in scene we care about; still checks objects not colliding (still bad!)
      - 10 objects with 100 triangles needs tens of thousands of comparisons
    - Better way: bounding volume
      - Sphere is simplest; not exact representation but only needs 50 comparisons
      - If two spheres intersect, check more closely (spend time only when needed)
      - Or test itself is enough if objects are far away enough and close inspection is unneeded
      - Too bad spheres aren’t always a good choice for objects; good fit is desirable
    - Bounding rectangle can be better, but intersection tests more complex
      - Simplest box is an “axis-aligned bounding box” or AABB
      - Oriented bounding box can be best, but even more complex for intersection tests (tradeoffs!)
    - Bounding shapes affect accuracy/speed, but they all use distance  $d = \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2 + (z_2 - z_1)^2}$ 
      - Ignore sqrt for speed!
    - Desirable Properties: inexpensive intersection tests (at least before restoring to expensive ones), tight fitting, inexpensive to compute BV, easy rotation/transformation, uses little memory
  - AABB
    - Simple to compute and straightforward!
    - Several ways to represent: Min/max extreme points (6 floats), min point and extends (6 floats or 3 floats + 3 half), center point and half-widths (same as above); last two are memory efficient
    - Computationally, min-extent is slowest and uses most operations
    - Intersection occurs if all axes overlap; if one axis is separated, no collision!

- Bounding sphere intersect if squared distance between centers  $\leq$  squared sum of radii
  - Computing bounding sphere itself isn't so simple; brute force takes  $O(n^5)$
  - Ritter uses iterative algorithm: start with sphere of 2 points, iteratively add points to grow (good for class)
    - Works quite well, but quality is sensitive to order of points added
  - Welzl uses computation geometry to achieve  $O(n)$  time, but complex implementation
- Processing BVs
  - Needs data structure
  - Organization depends on application
    - If we just need to check if spaceship hits asteroids, just check those (not asteroid-asteroid)
  - Advanced – what about hierarchies of bounding volumes?
    - If represented objects move, BV of bounding values have to update
  - Many bounding object types; spheres and AABB are less complex; maybe use sphere vs. plane (simple)

## Lecture 9:

- Representing Models
  - Typical model rendering: make instance, set model view/shader, draw, repeat...
  - More objects to render makes this approach unmanageable
  - Objects are also usually related to one another (car has wheels stuck to chassis, chassis relation to road/light post)
  - Use **scene graphs** (or trees): objects can be represented in logical way
    - Object relationships set in hierarchical relation to other objects (ex. wheels related to chassis, not world)
  - The graphs/trees may look like TRANSFORM  $\rightarrow$  CHASSIS  $\rightarrow$  (TRANS  $\rightarrow$  RFWHEEL, TRANS  $\rightarrow$  LFWHEEL, etc...)
  - TRS (translate, rotate, scale) transformation in each transform node
  - To make entire car move, only need to move top layer
  - Processing scene graph
    - Traverse using DFS (depth first!)
    - Transformations processed like stack as we descend into tree (with variations on exact representation)
      - Push/save current transformation on way down
      - Pop/restore previous transformation on way up
    - May also want to track other application/graphics state in this way
  - Elements of scene graph
    - More flexibility to keep transformations separate from geometry
    - But we could fatten scene by coming transform nodes; or even apply transformations to static geometry
    - Node types: root, transform, object (instance: transform+object, grouped); also includes...
      - Group (common parent)
      - Attribute (set graphics state, ex. texture)
      - Predicates/switches (render particular child of node based on rule or index)
    - Each node type defines some predefined behavior, maybe callbacks for application-specific functionality
  - Root node: parent of all nodes in scene graph
    - Global state about scene kept here (e.g. lighting)
    - Storing camera info does NOT go here; scene graph represents scene alone
  - Group node: groups children nodes/graphs
    - Usually just a tree structuring element, or often a base class for many other node types
  - Transformation node: applies transformation to all child nodes
    - Usually based on group node
    - Two versions in some scene graph APIs
      - Dynamic (transforms change over time); static (no change, can be flattened for slow processors)
  - Object node: stores object geometry
    - Simple scene graph could store references to everything needed to draw
    - Often actually stores pointers to data structures representing data (vertices, colors, normal, texture coordinates, textures, shaders, etc.)
  - Most scene graphs also compute BV for each node in tree
    - Bounding sphere for all group based nodes is common
      - Root, group, transform, etc.
    - Object nodes with geometry usually use tighter-fit BV (used for culling scene graph or collision detection)
    - Culling determines which parts of scene are outside of current view frustum
      - Only parts inside or partially inside frustum are drawn (partially intersecting ones are drawn too)
      - Involves comparing BVs of scene with six planes making up the view frustum



- View frustum planes:
- Object completely inside or intersecting one or more planes must be rendered; ignore the rest
  - Need plane equations of all six sides of frustum; extract planes directly from projection matrix!
  - Easiest to extract planes in clip space, where view volume (frustum) is cube around origin
    - Coordinates are still homogenous coordinates
  - After applying MV and projection matrix, vertices are in homogenous clip coordinates
  - Normalizing performs perspective division, where view volume at  $(-1,1), (-1,1), (-1,1)$ 
    - $pc = (xc, yc, zc, wc) = PMp$
    - $pcn = (xc/wc, yc/wc, zc/wc) = (x', y', z')$
  - But no access to values after perspective division (it's in hardware); but pc inside frustum if...
    - $Pc = (xc, yc, zc, wc) = PMp$
    - $-wc < xc, yc, \text{ and } zc < wc$
    - This means if  $-wc < xc$ , then pc is to the right of the left plane of the frustum
    - Recall  $p = (x, y, z, 1)$  and  $PM = [a_{11}, a_{12}, a_{13}, a_{14}; a_{21}, a_{22}, a_{23}, a_{24}; \dots; a_{41}, a_{42}, a_{43}, a_{44}]$
  - If point p evaluated to 0, point is on plane; if  $>0$ , positive/inside; if  $<0$ , negative/outside
  - For intersecting with sphere, we need to normalize for distance values to mean something
  - Frustum planes must be re-extracted when camera/objects in scene move (basically every frame)
  - Trivial to look up values for all 6 planes; as scene graph is traversed, BV is checked against frustum; if it falls outside the frustum, that branch of scene graph can be skipped

#### Lecture 9.5 (Physics in TA session):

- Particles, rigid body, cloth, water... (also deformable body, snow, smoke, fluid)
- We live in physical environment; realistic graphics is physical realism (how things move and look)
- Physics based animation
  - Particles, rigid body, deformable body, fluids, etc.
  - Real time vs. offline computation
- Physics based rendering
  - Global illumination, ray-tracing, path-tracing
  - Transparency, refraction, translucency (ios7), etc.
  - Depth of field
  - Motion blur
- Or just classical physics (no modern physics); things are getting to be real time with GPU power
- Particle Physics (trajectory of ball in angry birds)
  - For particle i:  $m_i$  (mass),  $x_i$  (position),  $v_i$  (velocity)
  - $f(t) = ma(t) = m(dv/dt) = m(d^2x/dt^2)$
  - Numerical Schemes:
    - Forward/explicit Euler method: use current time step force and velocity
    - Backward/implicit Euler method: use future time step force and velocity
    - Big difference when force dependent on position of other particles (when position independent, no diff)
    - In code:
      - To start, initiate position, velocity, and force for each particle object
      - Per frame, calculate force (ex. constant gravity), used to update velocity, used to update position; draw
- Rigid Body (collapse in angry birds)
  - Three translational ( $\chi$ ) and rotation ( $\theta$ ) degrees of freedom
  - In code:
    - To start, initiate center of mass, moment of inertia, rotation, velocity, angular velocity, and force per rigid body
    - Per frame, calculate force/torque (no collision), update velocity/angular velocity, update position/rotation; draw
- Cloth, hair, 1D string; spring-damper system (slingshot in angry birds)
  - Newton's law of motion:  $mass * acceleration = force$ ;  $ma = f$ ;  $m(d^2x/dt^2) = f$
  - System of second-order ordinary differential equations in time; net nodal force is  $f = s - \gamma * v + g$
  - $-\gamma * v$  is nodal drag/damping force,  $g$  is gravity,  $s$  is spring force

- Spring Force: net spring force at node  $i$  is the sum of forces due to springs connecting node  $i$  to neighboring nodes
  - $S(t) = \sum(S_{ij})$  for all  $j$  in neighbors; where  $S_{ij} = c_{ij}e_{ij} \cdot r_{ij} / |r_{ij}|$ ;
    - $r_{ij} = x_j - x_i$ , separation of nodes;  $|r_{ij}|$  is actual length;  $e_{ij} = |r_{ij}| - l_{ij}$ , deformation of spring
- A damped spring: parallel combination of spring and damper (Voigt model)
- In code:
- To start, init properties of nodes and spring-dampers
- Per frame, using forward euler: calculate sum of all forces, update velocity/position of nodes; draw everything
- Other forces
  - Friction force (stabilize the system): proportional to speed of node; friction factor
  - External force (manipulation): must solve pick problem; direct add to total force on picked node
- Advanced Deformable Object Simulation
  - Achieved by solving systems of partial differential equations
  - Continuum mechanics define equations
  - Apply advanced numerical methods to solve equations (Finite Element Method)
- Fluid and Smoke simulation
  - Navier-Stokes Equations; solved by applied mathematics
- Physics on the web?
  - How GPU shader-based simulation works
    - Fast, 10x improvement, no CPU resources used
    - GPGPU: general purpose computing on GPU (old way: programmable shaders; new way: CUDA/OpenCL)
    - WebGL? Supported on Chrome/Firefox by default; WebGL? Under development, but will be better
    - Shader:
      - Pipeline iterates at >30 fps; vertex and fragment shaders execute in parallel on many GPUs
    - Simulation shaders have two problems:
      - First, set of stat data need to be accessible/updatable at every time step
        - Solution: read/write all state data to two float point textures)
      - Update of data needs to be done locally in parallel fashion
        - Solution: most can simulations can be discretized spatially to a local update scheme
  - Cloth simulation algorithm
    - One of earliest physics simulation feat in graphics; 2D manifold embedded in 3D space, elastic material
    - Pseudo-physical model: internal forces (elastic, damping, curvature) and external (gravity, wind, drag)
  - Height-field water simulation and rendering of caustics
    - Represent fluid surface as 2D function  $h(x,y)$  or 2D array if discretized (simple, but can't do wave breaks)
    - Caustics: caused by redistribution of photons via refraction through curved surface
      - Ex. Light through water surface on pool bottom, or light through glass on table
      - Assume surface of caustic pattern is diffusive in all directions; generate photo intensity map on surface, then blend that with surface color texture for caustic pattern
    - Two-Pass Photon Map Generation
      - Project surface triangles to bottom via refraction; calculate total photons from surface triangle ( $area \cdot \cos(\theta)$ ,  $\theta$ =angle between surface normal and reversed incoming light direction)
      - Add up photon intensity from different projected triangles (turn on blending, set blending to add up fragment colors); photon intensity of projected surface triangle is total photons from triangle divided by area of projected triangle


## Lecture 10:

- Volume Rendering
  - We've only looked at surfaces in 3D space, not how values over 3D field might be rendered
  - Usually volume rendering deals with how to display a scalar field
    - Function that returns scalar value at particular location:  $scalar = f(x,y,z)$
  - Challenges: lot more data needed to represent volume, managing volume data is hard (do we store in file?!), how do we display it meaningfully?; here, we try to display internals of object too, not just surface
  - Data: could be sampling of physical process (like medical scan, physical simulation of some type, etc.)
    - Results in a 3D array of scalar sample values
    - One of 3D cells containing sample is a voxel, which represents a scalar value averaging that cell's value
    - Equally spaced 3D array of voxels is often a structured data set; otherwise, it's unstructured data set
  - Rendering
    - Direct volume rendering: use every voxel to produce an image
    - Iso-surface rendering: use implicit equation determines what's rendered



- Rendering iso-surfaces

- Implicit equation is  $c=f(x,y,z)$
- Extension of contours to 3D
- Finding “surface” that exists where the implicit function describing our data equals particular value
- Scalar values of this data field could represent things like temperature, pressure, density, etc.
- Want a surface that shows where value is true within data set (ex. where in 3D field is temp=275 deg?)
- Could project ray from eye/camera through every pixel on display and into data field
  - This requires being able to compute point along the line where value is equal to our target value
    - Ex. Where along the ray is the scalar value equal to value of interest
- Computationally expensive; no simple way to do this unless we use basic, known, quartic (like sphere)
- Also, volume data is usually discretized and not continuous; both of these are problems...
- Marching cubes is the best solution



- Form a surface mesh representing approximate value in question over discretized scalar field
- Representation of voxel with corners defined by scalar function  $f$ :



- Compare vertices of these cells to iso-surface value in question,  $c$ 
  - Each vertex is colored based on whether value is greater or less than that value
- With symmetry, that means 14 total variations




- Specific intersections along each edge can be found by interpolation
- Triangles used to tessellate portion of desired iso-surface mesh passing through the cell



- Each voxel contributes of eight of these cells
- When processing, move from one cell to next, row by row, then plane by plane (marching cubes)
- For each cell processed, the triangles are sent directly to rendering pipeline
  - Reasonable to collect triangles and build more efficient mesh
  - Ambiguities in the resulting mesh are ok, but not enough info in data to always resolve
- Iso-surfaces make rendering volumes manageable by really reducing amount of data to render
  - Only a fraction of cells get rendered

- Rendering direct methods


- Iso-surface is useful if you know value specified will show what you’re looking for (also pretty simple)
- But it only shows a single value, and the rest of the data is ignored
- To use contributions from all the data, need some mapping of scalar values to color and opacity
  - Otherwise, no idea how each value contributes
- Color/opacity selection is more art than science, but info on data can help make choices
- Peaks can tell us which elements are interesting to call out (ex. Bone, soft tissue, etc.)



- 
- Adjust opacity of any color to call out particular structure within the volume
  - Usually user sets values interactive
- This is known as Transfer Function




- 
- Splatting: simple method where shape assigned to all voxels
  - Exact shape to use for a splat is undergoing research (often ellipsoids used)
  - Shapes projected/composited back-to-front on viewing plane so transparency/blending is ok
  - Order defined by data grid orientation to view point; spatial data structure (oct-tree) can sort
- Ray casting: similar to implicit method of iso-surface rendering in that we project ray through volume
  - Values along ray are tri-linearly interpolated to color and opacity values



- 
- Each sample is lit normally using gradient of voxel with respect to light source and viewer
- Finally, samples composited back-to-front order to obtain final pixel value
- Very high quality results! GPUs can be used to speed this technique up a lot
- Textures:
  - OpenGL supports 3D textures; each slice/plane of a data volume is stored in texture map
    - Colors/opacities baked onto images
  - Applying 3D texture coordinates to geometry allows arbitrary slices of data to be rendered
  - Sampled values are tri-linearly interpolated from texture data


## Lecture 11:

- Parallel Rendering
  - Sometimes one computer isn't enough: not enough pixels OR not enough polygons
  - Not enough pixels
    - Even high-res displays aren't enough (HDTV is 2 megapixels); tile displays together as a "power wall"
    - Tiled displays can reach 150 MP+ (ex. SAGE, a library used to manage these walls developed by UIC)
  - Not enough polygons
    - Parallelizing increases throughput; but how we parallelize depends on problem
    - All graphics parallelizing breaks down into either: sort first, sort middle, or sort last
    - Sort middle: how all GPUs work nowadays
      - Any number of geometry process (G) and/or fragment rasterizers (R)




- Each rasterizer associated with part of display
  - Primitives sorted to rasterizing corresponding to projected area of primitive
  - Load balances fairly well on GPU; difficult at application level
- Sort last: geometry and rasterization handled by single unit

- Load balance across all units



- Good load balancing for rendering, but we'll have to composite all pieces back together
    - Compositing is a problem...
    - Each rendering system can potentially render to entire display
    - To composite result requires depth information
    - Compositing reads entire color and depth buffer and sends over bus or network
      - All buffers are combined into final image, so fast network is required
      - Load balancing is great, but speed can be an issue
        - Readback, network, depth processing, upload
- Sort first: objects ordered to render that handles the part of the display it will be projected to

- Hard to load balance, but compositing is faster




- Trick is to quickly determine where objects will land onscreen
- Must estimate how long it will take to render objects in any scheme
  - One way around this is to adjust screen partitions
- Videos

## Lecture 12:


- Shadows
  - Important to the real-world light experience; helps understand shape and spatial relationships between objects
  - Challenging in real-time graphics
    - Algorithms with simple concepts are too slow in real-time
    - Even today, no technique works well in all situations (best approach depends on constraints on situation)



- Def: “area that’s not or partially illuminated due to light being intercepted by opaque object between area/light”
  - Shadow graphics definition: “region of space for which at least one point of light source is occluded”
  - Assumptions: only direct illumination considered (no bouncing light); occluders assumed as opaque
- Shadows in graphics
  - Completely general algorithms extending beyond opaque objects and direct illumination are still very much beyond current hardware capabilities for real-time rendering
  - Accurate shadows are among most important unsolved problems in graphics; we’ll use simple definition
- Shadow components
  - Point p on surface can be one of the following with respect to an area light source
    - Entire light source blocked by scene → p within umbra and in shadow
    - Light source partially blocked by scene → p within penumbra and partially in shadow
    - Light source not blocked by scene → p is lit and not in shadow



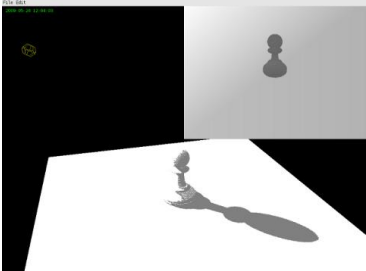
- Basically, trace a ray from light source l to point p on a surface
  - If intersection between l and p, point p is in shadow
  - Object intersected by ray is the “occlude” (or the “blocker” or “shadow caster”)
  - Object point p belongs to is the receiver
- To simplify more, we’ll only use point lights to avoid difficult integration needed for area light sources
  - Only using point light sources restricts results to hard shadows only; why?
  - Soft shadows can be achieved using approximation techniques that give appearance of sampling an area light source (ex. Percentage closer filtering); we’re not going to discuss these
- Techniques for hard shadows
  - Shadow mapping (projective & depth)
  - Shadow volumes (not going to discuss)
  - All variations are tradeoff between efficiency and accuracy
  - “Hard” is used because the result is binary, either point is light or it’s not (either in shadow or not)
  - Remember there’s no point light source in the real world - not even a light bulb!
    - Approximation to simplify computation in order to achieve interactive frame rates
- Planar projective shadows
  - Projection of shadow onto a planar (flat) surface only!
  - Lots of problems, not really used ever
- Projective shadow texture
  - Variation of projective textures that’s not limited to planar receiver surfaces
    - Think of projector at light source casting an image of the shadow onto receivers




- - Texture buffer is cleared to white
  - Occluders are rendered in black into buffer (a single texture lookup tells us if receiver is in shadow or not)
  - To render into texture, setup camera in light space, NOT camera space
  - Good use for lookAt function to setup appropriate projection matrix
    - Position camera at light source location, in direction of the shadow (light)
  - Results in transformation matrix that transforms vertex from world space to light space
  - Need projection matrix to transform our points into “light clip space”
  - Use orthographic or perspective transformation
    - Common to use orthographic projection for infinite light sources (ex. Sun)

- Point light source would use perspective transform
- $d$  is the distance from light source to projection plane
  - $M_{L,V} = [d \ 0 \ 0 \ 0; 0 \ d \ 0 \ 0; 0 \ 0 \ d \ 0; 0 \ 0 \ 0 \ -1]$
- We're in clip space: everything needs to map to  $[-1,1]$ , so  $d$  is typically 1, which projects to  $z=-1$
- To keep  $x, y$  in same range requires small change to projection matrix with additional scaling values;  $w, h$  are size of our shadow texture buffer
  - $M_{L,P} = [w \ 0 \ 0 \ 0; 0 \ h \ 0 \ 0; 0 \ 0 \ 1 \ 0; 0 \ 0 \ 0 \ -1]$
- Not useful for this technique, but for others: maintain depth info rather than projection everything to  $z=-1$ 
  - Results in familiar perspective projection matrix; you can use `Perspective()` function
  - Similar to results with a Kinect
- In vertex shader when rendering receivers, shadow texture's texture coordinates are computed as  $v^s$ 
  - Need to adjust light clip values  $[-1,1]$  to conventional text coordinates  $[0,1]$ 
    - $M_t = [0.5 \ 0 \ 0 \ 0.5; 0 \ 0.5 \ 0 \ 0.5; 0 \ 0 \ 0.5 \ 0.5; 0 \ 0 \ 0 \ 1]$
    - Entire projection baked into  $M_s = M_t * M_{L,P} * M_{L,V}$  and  $V^s = M_s * v$
- Values outside of  $[0,1]$  not shadowed (checked for in fragment shader)
- Issues, both good and bad
  - Must separate occluders and receivers; requires shadow texture per blocker; no self shadowing
  - Can achieve a form of soft shadow by filtering texture map
    - This "light attenuation map" bakes shadows of static lights/occluders/receivers to scene
  - Overall, very simple technique

## • Shadow Mapping

- More general form of projective texture shadows
  - No need to separate occluders from receivers; can handle self-shadowing
  - Again, render from position of light source
    - But render entire scene; every point rendered is implicitly lit (anything else is in shadow)
    - Determining whether 3D position in shadow or not needs checking whether lit in shadow map or not
- 
- Simple in theory, hard to do nicely in practice
  - Image spaced based, so artifacts are necessary; shadow map resolution can result in jagged shadow edges
  - We're only interested in depth map created by rendering from light's position
  - Each fragment's position  $p$  is transformed into light clip space, like before
  - The  $x, y$  components index into the depth map;  $z$  value is distance from light source
  - Remember, actual  $x, y$  values need to be scaled like projective texture was to  $[0,1]$ , or  $p^s$
  - $Z$  value is compared to depth value of point under consideration
    - If point in light clip space has depth value  $>$  point's value in shadow map, the fragment is hidden or in shadow; otherwise, it's lit and shaded normally
  - Since we've been producing maps via projection and rendering...does this only work for spot (directional) lights?
    - No, typical solution involves rendering 6 frustums around light source, much like a cube map
      - Spherical maps more efficient in terms of pixel/rendering geometry, but cube maps are simple
      - Some have suggested to compromise with tetrahedrons
  - Bigger issue: sampling shadow map buffer during depth comparison
    - Shadow map has limited precision to represent depth values
    - Shadow map has sampled scene at different resolution than camera
    - Rarely does the point in eye space projected into light space correspond to exact depth map sample
    - Leads to light "leaks" due to z-fighting
    - Occurs mostly when receiver is tilted and discretization of depth values result in incorrect comparison
  - Solution: introduce a "bias", a value which pushes the depth values slightly away from the light source
    - Problem is that there's no rule of thumb for deciding what the bias value should be (need to experiment!)
    - But there's a built-in OpenGL mechanism for applying bias, called `glPolygonOffset()`
  - From top, left to right: view from camera, view from light, shadow map from light, shadow map from camera, scene depth, final image after compare



- 
- Of course this assumes we can render to texture; in depth map case, we're rendering depth values to a texture
  - This is done using OpenGL frame buffer object (or FBO)
- In code:
- On setup, create appropriate texture
  - Notice the compare parameters which control how depth comparison will operate
- Next, create FBO and associate with texture we just defined; finally, render to the FBO

### Lecture 13:


- Curves
  - So far only done flat surfaces
    - Even sphere was approximated by repeatedly sub-dividing into flat sided object (tetrahedron)
    - Convenient because OpenGL designed to render flat objects (AKA triangles) very efficiently
  - New ways to model curves and surfaces
    - Implementation involves flat primitives (lines, triangles) but with description as curves
- Representations
  - Explicit: familiar, but has problems for us
  - Implicit: also familiar, but no analytic form we can use
  - Parametric: less familiar, but most useful for our purposes
- Explicit Representation
  - $y=f(x)$  form with one independent, one dependent variable
  - No guarantee this form exists for a given "curve"
    - ex. Line represented as  $y=mx+b$ , but can't represent vertical line
    - Likewise, many curves have no explicit form (ex. A loop that looks like gamma)
  - Circle is even more obvious of a problem; can only represent half of circle
    - We'd need to do  $y=\sqrt{r^2-x^2}$  AND  $y=-\sqrt{r^2-x^2}$ , iff  $0 \leq |x| \leq r$
  - 3D has similar problems as 2D
    - Curves need 2 equations (dependent variables):  $y=f(x)$ ,  $z=g(x)$
    - Surfaces require two independent variables:  $z=f(x,y)$
    - Some curves not represented in this form either
      - Lines can't be described if it exists on any plane of constant  $x$ , when defined in terms of  $x$ 
        - Kind of like vertical line problem
      - Spheres can't be described due to same ambiguities as 2D case without constraints
        - $z=f(x,y)$  can generate 0,1, or 2 points on sphere
      - Too many corner cases to be useful for more complex curves and surfaces (patches)
  - Implicit Representation
    - Curves represented using form  $f(x,y)=0$
    - Line has usual  $ax+by+c=0$  form; circles look like  $x^2+y^2-r^2=0$ 
      - No real representational limitations here!
    - Except implicit form is more useful for testing membership
    - Is point on line, curve, surface, or not? Collision detection!
    - Generally no convenient way to analytically determine  $x$  given a  $y$  value; this limits usefulness in rendering
  - Same for 3D: we can represent lines, curves, surfaces
    - Lines are  $ax+by+cz+d=0$ , surfaces are  $f(x,y,z)=0$
    - Curves are trickier with intersection of two 3D surfaces:  $f(x,y,z)=0$  and  $g(x,y,z)=0$
  - While almost any curve/surface we use has implicit representation, extracting points along or on them is difficult!
    - We need those points to draw the curve!

- Parametric Representation

- Represents each dimensional point on curve with respect to a single independent variable  $u$ 
  - $x=x(u), y=y(u), z=z(u)$ ; only use  $z$  for 3D representation
  - Varying  $u$  allows us to generate points that sweep out the curve
- Convenient for rendering, but not necessarily fast
- Surfaces are similar:  $x=x(u,v), y=y(u,v), z=z(u,v)$ ; these can all go in a vector  $p(u,v)$ 
  - Varying  $u, v$  can sweep out a surface
  - Taking cross product of tangent vectors at a point gets us the normal of the surface
    - $n = dp/du \times dp/dv$
- Curves represented as polynomials in terms of  $u$  and surfaces in terms of  $u, v$  are convenient to use; we get:
  - Local control of their shape
  - Smoothness and continuity control between curves/shapes
  - Ability to evaluate derivatives
  - Behavioral stability
  - Ease of rendering


- Example curve

- $p(u)=[x(u), y(u), z(u)]$
- Polynomial representation of degree  $n$  is (not dimension):  $p(u) = \sum_{k=0}^n (u^k * c_k)$ 
  - Each of  $n+1$   $c_k$ 's has independent coefficients, like  $c_k=[c_{xk}, c_{yk}, c_{zk}]$




- Example surface

- $p(u,v)=[x(u,v), y(u,v), z(u,v)]$
- Polynomial representation of degree  $n$  is  $p(uv) = \sum_{i=0}^n (\sum_{j=0}^m (c_{ij} * u^i * v^j))$
- Again,  $m+1$  and  $m+1$   $c_k$ 's have independent coefficients
- Generally,  $n=m$  generates a square shaped surface patch




- Choosing degree: we can use whatever degree we want

- But high degree results in complex curve that's hard to control
- Low degree might be too simple to represent what we need
- Solution: use collection of simpler curves to represent more complex shapes
- In graphics, most common degree is 3, resulting in cubic polynomials (we'll almost always use this)
- Degree of 3 keeps shape control local and lets us manage joints between curve segments more easily
  - Degree of 3:  $p(u)=c_0 + c_1*u + c_2*u^2 + c_3*u^3 = \sum_{k=0}^3 (c_k * u^k) = u^T * c$
  - $c = [c_0, c_1, c_2, c_3]$ ;  $u=[1, u, u^2, u^3]$
- We have to solve for values of  $c$ 
  - 12 equations in 12 unknowns for degree of 3
  - $X, y, z$  are independent, so we group the problem into 3 sets of 4 equations with 4 unknowns
  - Use set of control point data to help solve for unknowns, lets us generate points along curve



- Cubic interpolating polynomial

- Not used often in CG
- There are other curve types with beneficial properties for control and rendering
- For now, have four control points in 3D
- Seek coefficients  $c$  such that polynomial  $p(u)=u^T c$  passes through control points




- Control points defined by  $p_k = [x_k, y_k, z_k]$
- Also choose values of  $u$  at each  $p$  to perform interpolation; these are convenient:  $u=0, 1/3, 2/3, 1$
- This gives four conditions, where
  - $p_0 = p(0) = c_0$
  - $p_1 = p(1/3) = c_0 + 1/3 c_1 + (1/3)^2 c_2 + (1/3)^3 c_3$
  - and so on, until  $p_3 = p(1) = c_0 + c_1 + c_2 + c_3$
- Thus we get  $p = Ac$ , where
  - $p = [p_0, p_1, p_2, p_3]$
  - $A = [1 \ 0 \ 0 \ 0; 1 \ 1/3 \ (1/3)^2 \ (1/3)^3; \dots; 1 \ 1 \ 1 \ 1]$
  - Inverting  $A$  gives “interpolating geometry matrix” and the coefficients  $c = M_I^{-1} p$
- Now we can evaluate points along curve
- To continue our curve, define next segment with  $p_3, p_4, p_5$ , and  $p_6$ ; achieves continuity!
  - Does not ensure the same derivative at point  $p_3$ , where curves join
  - May or may not be a problem
- We mentioned we desired ability to build up complex curves using simpler ones
  - Behavior of how curves join says plenty about how we achieve this!


#### Lecture 14:

- Curves
- Parametric: Cubic Blending Functions
  - Slightly different way of looking at this interpolation process; it allows us to see what exactly is going on
  - Let's substitute interpolating coefficients into polynomial itself (generalize things a bit)
  - $p(u) = u^T c = u^T M_I^{-1} p$ , or  $p(u) = b(u)^T$ , where  $b(u) = M_I^{-1} u$
  - $b(u) = [b_0(u), b_1(u), b_2(u), b_3(u)]$
  - Expressing  $p(u)$  in terms of cubic blending polynomials gives us:
    - $p(u) = b_0(u) p_0 + \dots + b_3(u) p_3$
    - $b_0(u) = -9/2 (u-1/3)(u-2/3)(u-1)$
    - $b_1(u) = -27/2 u(u-2/3)(u-1)$
    - $b_2(u) = -27/2 u(u-1/3)(u-1)$
    - $b_3(u) = 0/2 u(u-1/3)(u-2/3)$
  - Lets us see how each blending equation factors into interpolation
- Each one isn't particularly smooth (since we're having interpolation pass through each control point)
- Higher degree polynomials have more pronounced swings
- Remember there's no way to enforce derivatives at endpoints; makes this form limited
- Parametric: Hermite Form
  - We can form curve/surface using cubic interpolating polynomial, but there are issues
  - Hermit form allows additional control over derivatives at ends of the curve
  - Here, only consider control points  $p_0$  and  $p_3$ , which from previous example, has first two conditions
    - $p_0 = p(0) = c_0$
    - $p_3 = p(1) = c_0 + c_1 + c_2 + c_3$
  - We get other two conditions if we assume the derivatives at  $u=0$  and  $u=1$  are known
    - $P'(u) = c_1 + 2uc_2 + 3u^2 c_3$
    - $p_0' = p'(0) = c_1$
    - $p_3' = p'(1) = c_1 + 2c_2 + 3c_3$

- matrix form:  $q=[p_0, p_1, p_2, p_3]=[1 \ 0 \ 0 \ 0; 1 \ 1 \ 1 \ 1; 0 \ 1 \ 0 \ 0; 0 \ 1 \ 2 \ 3]$




- Solve for  $c$  to get  $c=M_H^{-1}q$ ,  $M_H=[1 \ 0 \ 0 \ 0; 0 \ 0 \ 1 \ 0; -3 \ 3 \ -2 \ -1; 2 \ -2 \ 1 \ 1]$ 
    - $M_H$  is the Hermite geometry matrix
  - Again, we get resulting polynomials
  - Blending functions can be used in same way as before for cubic interpolating polynomial
  - Hold derivative to be the same across curve segments (at the join) to get continuity



- Curves connecting at endpoints have  $C^0$  parametric continuity
  - If derivatives also match, we get  $C^1$  parametric continuity
  - If derivatives are proportional to each other, we get  $G^1$  geometric continuity


#### • Parametric: Bezier Form

- Can't really compare cubic interpolating polynomial to Hermite
  - Both are cubic in degree; but they don't use same data (control points)
- Bezier Form: use all four control points of cubic interpolating polynomial to approximate Hermite curve
  - Named after Pierre Bezier, who worked for Renault in France (1960s)
  - Again, use endpoints  $p_0$  and  $p_3$ , insisting interpolation passes through these values
    - $p_0 = p(0) = c_0$
    - $p_3 = p(1) = c_0 + c_1 + c_2 + c_3$
  - Bezier uses  $p_1$  and  $p_2$  to approximate tangents at  $u=0$  and  $u=1$ , instead of using them for interpolation




- Approximating tangent results in some conditions...
- Get four equations with four unknowns
  - Solving for  $c$ ,  $c=M_B^{-1}p$ , where  $M_B=[1 \ 0 \ 0 \ 0; -3 \ -3 \ 0 \ 0; 3 \ -6 \ 3 \ 0; -1 \ 3 \ -3 \ 1]$
  - $M_B$  is Bezier geometry matrix
- Cubic polynomial is then  $p(u) = u^T M_B^{-1} p$ 
  - Exact same manner as cubic interpolating polynomial seen earlier
  - If control points are overlapped, you should see that we still have  $C^0$  continuity at join
  - We do not have  $C^1$  continuity like with Hermite due to different approximations for tangent


- Blending functions:  $p(u)=b(u)^T p$ , where  $b(u)=M_B^{-1}u=[(1-u)^3, 3u(1-u)^2, 3u^2(1-u), u^3]$



- Zero values only at end of interval; this ensures smooth interpolation over interval  $[0,1]$
- Also see that while  $0 < u < 1$ 
  - Blending functions are also  $b(u) < 1$ ; this condition is called a "convex sum"
  - Implies that curve will be contained within convex hull of control points (good interactive design)



- Parametric: Cubic B-Splines
  - Bezier curves only achieve  $C^0$  continuity (mathematically speaking)
  - We can achieve  $C^1$  by matching tangents
  - Relax condition that interpolation must pass through control points to achieve  $C^2$  continuity using cubic B-spline



- Previously, we varied  $u$  from 0 to 1, the curve spanned over all four control points
    - Instead, consider spanning over only middle two control points
  - Matching conditions at  $p(0)$  with  $q(1)$  achieves  $C^2$  continuity by solving for  $M$
  - $p(u) = u^T M^* p$ , where  $p = [p_{i-2}, p_{i-1}, p_i, p_{i+1}]$
  - $q(u) = u^T M^* q$ , where  $q = [p_{i-3}, p_{i-2}, p_{i-1}, p_i]$
- 
- Use symmetric approximations for tangent at joint point (like  $p_{i-2}$ ,  $p_{i-1}$ , and  $p_i$ ) to get
    - $p(0) = q(1) = 1/6 * (p_{i-2} + 4p_{i-1} + p_i) = c_0$
    - $p'(0) = q'(1) = 1/2(p_i - p_{i-2}) = c_1$
  - Do same conditions at  $p(1)$ , sliding down to next set of control points
    - $p(1) = 1/6 * (p_{i-1} + 4p_i + p_{i+1}) = c_0 + c_1 + c_2 + c_3$
    - $p'(1) = 1/2(p_{i+1} - p_{i-1}) = c_1 + 2c_2 + 3c_3$
  - Now we have four equations for coefficients of  $c$ , so we can solve for  $M$
  - B-spline geometry matrix,  $M_S = 1/6 * [1 \ 4 \ 1 \ 0; -3 \ 0 \ 3 \ 0; 3 \ -6 \ 3 \ 0; -1 \ 3 \ -3 \ 1]$
  - We now have  $C^2$  continuity at joins, but at a least of 3 times the amount of work!
    - Had to interpolate between each set of control points
  - $C^2$  continuity not only connects segments, but matches tangents and curvature
    - Very useful properties when modeling real world materials
    - Difficult to use, since curve does not pass through any of control points – not intuitive